

STYX: Stream Processing with Trustworthy Cloud-based Execution

Julian Stephen, Savvas Savvides, Vinaitheerthan Sundaram, Masoud Saeida Ardekani, Patrick Eugster

October 6, 2016

Purdue University

Table of contents

1. Overview
2. Ensuring confidentiality in the cloud
3. Challenges in encrypted stream processing
4. Architecture
5. STYX abstractions and key update
6. Evaluation
7. Related work and conclusion

Overview

Compute clouds

- Data analytics platforms
- Cost-efficiency, 'on-demand' compute, low infrastructure setup cost

IoT

- 26 billion smart devices connected to a network by 2020
- Fine-grained user behavior tracking to capture, personalize and/or monetize user experience

Stream processing

- Analytics on real-time streaming data (continuous queries)
- Many systems over the last few years - Apache Storm, Apache Spark, Apache Flink, Apache Samza, Amazon Kinesis




Vulnerabilities - 1




Vulnerabilities - 1




Vulnerabilities - 1



Vulnerabilities - 1



Vulnerabilities - 1



Real problems

Xen Security Advisory CVE-2014-7188 / XSA-108

IMPACT

=====

A buggy or malicious HVM guest can crash the host or read data relating to other guests or the hypervisor itself.

Public release.

Ensuring confidentiality in the cloud


Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data

$$A \longrightarrow f \longrightarrow f(A)$$


Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data



Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data




Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data

$$\begin{array}{ccccc} A & \longrightarrow & f & \longrightarrow & f(A) \\ \downarrow k & & \downarrow & & \\ E(A, k) & \longrightarrow & f' & \longrightarrow & f'(E(A, k)) \end{array}$$

Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data



Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data
- Prohibitive overhead

Confidentiality in the cloud

Fully homomorphic encryption (FHE)

- Allows arbitrary computation on encrypted data
- Prohibitive overhead

Partially homomorphic encryption (PHE)

- Allows certain operations to be performed over encrypted text
- AHE: $D(E(x1) \psi E(x2)) = x1 + x2$
- AHE, MHE, OPE, DET

Conjecture

Many data analytics jobs can be performed securely using a combination of partially homomorphic encryption schemes


Vulnerabilities - 3




Vulnerabilities - 3




Vulnerabilities - 3



Vulnerabilities - 3




Vulnerabilities - 3



Challenges in encrypted stream processing

Challenges in encrypted stream processing

- Programmer effort
 - Need to identify encryption scheme for each input data stream, perform cryptographic equivalent of required operation



- `if (Stream1.f1 < 100)`
 `return; else ...`
- `sum = sum + Stream2.f3;`


Challenges in encrypted stream processing

- Key change
 - PHE requires all tuples in an aggregate function to be encrypted with same key



Challenges in encrypted stream processing

- Deployment optimizations
 - Deployment parameters specified for plaintext data may not be optimal when computation happens on encrypted data



Challenges in encrypted stream processing

- Programmer effort
 - Need to identify encryption scheme for each input data stream, perform cryptographic equivalent of required operation
- Key change
 - PHE requires all tuples in an aggregate function to be encrypted with same key
- Deployment optimizations
 - Deployment parameters specified for plaintext data may not be optimal when computation happens on encrypted data

Challenges in encrypted stream processing

- Programmer effort
 - Need to identify encryption scheme for each input data stream, perform cryptographic equivalent of required operation
- Key change
 - PHE requires all tuples in an aggregate function to be encrypted with same key
- Deployment optimizations
 - Deployment parameters specified for plaintext data may not be optimal when computation happens on encrypted data
- Limitations of PHE
 - PHE may not support a sequence of operations requiring trusted nodes to perform remaining computation


Challenges in encrypted stream processing

- Programmer effort
 - Need to identify encryption scheme for each input data stream, perform cryptographic equivalent of required operation
- Key change
 - PHE requires all tuples in an aggregate function to be encrypted with same key
- Deployment optimizations
 - Deployment parameters specified for plaintext data may not be optimal when computation happens on encrypted data
- Limitations of PHE
 - PHE may not support a sequence of operations requiring trusted nodes to perform remaining computation
- Constants and initialization
 - Variables must be initialized using the encrypted value of the initialization constant

Challenges in encrypted stream processing

- Programmer effort
 - Need to identify encryption scheme for each input data stream, perform cryptographic equivalent of required operation
- Key change
 - PHE requires all tuples in an aggregate function to be encrypted with same key
- Deployment optimizations
 - Deployment parameters specified for plaintext data may not be optimal when computation happens on encrypted data
- Limitations of PHE
 - PHE may not support a sequence of operations requiring trusted nodes to perform remaining computation
- Constants and initialization
 - Variables must be initialized using the encrypted value of the initialization constant

Architecture



Execution flow

- User submits program written using system (STYX) API
- Homomorphism analysis identifies crypto systems required to execute the graph
- Analytical model identifies deployment profile
- Scheduler assigns tasks to nodes
- Runtime executes tasks

STYX abstractions and key update

Group sum in a sliding window

```
1  /** Track sum of values per group per time slot */
2  public class SlotBasedSum<T> {
3      ...
4      public void updateSum(T group, int slot,
                           SecField val) {
5          SecField[] sums = objGroupSum.get(group);
6          if (sums == null) {
7              sums = new SecField[this.numSlots];
8              init(sums, val);
9              objGroupSum.put(obj, sums);
10         }
11         sums[slot] = SecureOper
12             .add(sums[slot], val);
13     }
14 }
```

Group sum in a sliding window

```
1  /** Track sum of values per group per time slot */
2  public class SlotBasedSum<T> {
3      ...
4      public void updateSum(T group, int slot,
                           SecField val) {
5          SecField[] sums = objGroupSum.get(group);
6          if (sums == null) {
7              sums = new SecField[this.numSlots];
8              init(sums, val);
9              objGroupSum.put(obj, sums);
10         }
11         sums[slot] = SecureOper
12             .add(sums[slot], val);
13     }
14 }
```

Group sum in a sliding window (Storm)

```
1 public class SlotBasedSum<T> {
2     BigInteger publicKey = readPubKey();
3     public void updateSum(T group, int slot,
4                           BigInteger value) {
5         BigInteger[] sums = objGroupSum.get(group);
6         if (sums == null) {
7             sums = new BigInteger[this.numSlots];
8             init(sums, "AHE");
9             objGroupSum.put(group, sums);
10        }
11        sums[slot] = sums[slot].multiply(value)
12            .mod(publicKey.multiply(publicKey));
13    }
14 }
```

Group sum in a sliding window (Storm)

```
1 public class SlotBasedSum<T> {
2     BigInteger publicKey = readPubKey();
3     public void updateSum(T group, int slot,
4         BigInteger value) {
5         BigInteger[] sums = objGroupSum.get(group);
6         if (sums == null) {
7             sums = new BigInteger[this.numSlots];
8             init(sums, "AHE");
9             objGroupSum.put(group, sums);
10        }
11        sums[slot] = sums[slot].multiply(value)
12            .mod(publicKey.multiply(publicKey));
13    }
14 }
```

Group sum in a sliding window (Storm)


```
1 public class SlotBasedSum<T> {
2     BigInteger publicKey = readPubKey();
3     public void updateSum(T group, int slot,
4                           BigInteger value) {
5         BigInteger[] sums = objGroupSum.get(group);
6         if (sums == null) {
7             sums = new BigInteger[this.numSlots];
8             init(sums, "AHE");
9             objGroupSum.put(group, sums);
10        }
11        sums[slot] = sums[slot].multiply(value)
12            .mod(publicKey.multiply(publicKey));
13    }
14 }
```


Key change

Challenges

- Functions that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Problem




Key change

Challenges

- Functions that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Problem



Key change

Challenges

- Functions that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Problem




Key change

Challenges

- Functions that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Problem




Key change

Challenges

- Continuous queries that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Solution




Key change

Challenges

- Continuous queries that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Solution




Key change

Challenges

- Continuous queries that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Solution




Key change

Challenges

- Continuous queries that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output

Solution




Key change

Challenges

- Continuous queries that aggregates data over a sliding window makes it impossible to change the encryption key without disrupting output


Solution



Evaluation


IoT Bench

- Smart meter data over a 24 hour time period at the rate of 1 reading per minute from 443 unique homes, totaling 637526 records
- 4 ec2 m3.large node
- Each tuple comprises of timestamp, meter id and meter reading



Performance when keys change

- New york taxi route data (10G)
- Application finds the top 10 most frequent routes during the last 30 minutes of taxi servicing
- 9 ec2 m3.large node



Related work and conclusion

Prior work on encrypted computation

- [Popa et al.; SOSP'11];
- [Gentry.; STOC'09];
- [Stephen et al.; VLDB'13];

Other approaches

- Using secure processors (e.g., Intel SGX)

Conclusion

- Confidentiality breaches pose a serious threat to adoption and utilization of cloud resources
- PHE has proven to be effective for various batch workloads
- STYX leverage PHE for stream data analytics in the cloud
- Makes it easier for programmers to use PHE
- Automatically translates the program and optimizes deployment parameters

Questions